



Introduction To Algorithms

Week 8 - Final Synthesis & Assessment:
"Explaining the Solution You Submit"

1



Terms to Remember

Term	Brief Meaning
Practical consideration	Real-world issue that affects use
Explainability	Ability to explain how or why a result happened
Global explanation	Explains overall system behavior
Local explanation	Explains one specific output or decision
LIME	Local explainability technique
Bias	Systematic unfair skew or distortion
Discrimination	Harmful unequal treatment
Privacy	Protection of sensitive information
Traceability	Ability to follow evidence and decisions
Inconclusive evidence	Evidence that does not fully prove the claim
Misguided evidence	Evidence that points attention in the wrong direction
Unfair outcome	Result that harms or disadvantages unfairly
CRISP-DM	Data mining process lifecycle
Black swan event	Rare high-impact event hard to predict beforehand

2



What We Will Use Today

- ▶ Explainability
- ▶ Evidence quality
- ▶ Assumptions and limitations
- ▶ Ethical risk recognition
- ▶ Bias and privacy cautions
- ▶ Cost, time, and accuracy tradeoffs
- ▶ Black Swan Event awareness
- ▶ FINAL: Part 1 and Part 2 preparation
- ▶ Optional bonus-credit / enrichment application

3




What We Will Explore in Other Courses

- ▶ Formal AI ethics
- ▶ Applied model explainability
- ▶ Larger-scale data validation
- ▶ Security and privacy governance
- ▶ Production monitoring
- ▶ Advanced ML model evaluation
- ▶ Responsible AI system design



4

Review Coding Lab

5



Review: What Lab 07 Taught Us

Last week - A result depended on choices:

- data representation
- similarity rule
- ranking or grouping method
- assumptions
- limitations
- AI/data connection

6

Southwest Tech

Today's Question

After an algorithm gives an answer, what still needs to be explained?

7

Southwest Tech

Examine the Output with Good Questions

The diagram shows a central box labeled "Algorithm Output" with a bar chart icon. It contains the text "Result: Recommended Option A" and "Score: 86". Four arrows point to this central box from surrounding boxes:

- Evidence:** What supports this result? (Green box with magnifying glass icon)
- Assumptions:** What did we assume? (Blue box with cloud icon)
- Limits:** Where might this fail? (Purple box with warning triangle icon)
- Impact:** Who or what is affected? (Orange box with group of people icon)

Image Slide

8

Southwest Tech

Final Week is Synthesis

No Lab!

It is:

- ▶ final synthesis
- ▶ practical-risk discussion
- ▶ explanation practice
- ▶ final assessment preparation & exam

9

What Counts as Success Today

- explain evidence
- name assumptions
- describe tradeoffs
- identify limitations
- discuss responsible AI/tool use
- prepare for the final explanation defense

10

Southwest Tech

Textbook Review

Practical Considerations

Chapter 16 asks:
"What happens when algorithms meet real-world use?"

The reading covers:

- ▶ - unexpected behavior
- ▶ - explainability
- ▶ - ethics
- ▶ - bias
- ▶ - privacy
- ▶ - when algorithms fit
- ▶ - limits of prediction

11

Southwest Tech

Where Algorithmic Code Meets Real-World Constraints

The diagram shows a central box labeled "Algorithmic Solution" with a code icon. It is surrounded by five boxes connected by arrows:

- Users:** People with goals, needs, and behaviors.
- Data:** Real-world data with quality, biases, and gaps.
- Policy:** Rules, regulations, and ethical boundaries.
- Unexpected Inputs:** Edge cases, errors, and new situations.

At the bottom, a box with a globe icon states: "Working code still operates inside a larger situation."

Practical Considerations

12

Textbook Review

Southwest Tech

Expecting The Unexpected

Algorithmic solutions face challenges:

- ▶ unusual inputs
- ▶ missing data
- ▶ changing behavior
- ▶ misuse
- ▶ unclear goals
- ▶ real-world consequences

13

Textbook Review

Southwest Tech

Cautionary Example: Tay

Tay shows that systems can fail when:

- ▶ public input is uncontrolled
- ▶ learning behavior is not bounded
- ▶ safeguards are weak
- ▶ harmful patterns are amplified

14

Textbook Review

Southwest Tech

Explainability

Explainability asks:

- ▶ Why did the system produce this output?
- ▶ What evidence supports it?
- ▶ What inputs mattered?
- ▶ What assumptions shaped it?
- ▶ What should we not conclude?

15

Textbook Review

Southwest Tech

Global And Local Strategies

Global explanation:

- ▶ explains overall behavior

Local explanation:

- ▶ explains one specific output or decision

Both can matter.

16

Southwest Tech

Global vs Local Explanation

Global View

How the system usually behaves

Inputs: User Data, Usage Data, Context Data, History Data

Rules: Rule 1, Rule 2, Rule 3, ...

Outputs: Option A, Option B, Option C, ...

How the system usually behaves

Local View

Why did this one result happen?

Input (this case): User: Alex

Rule Applied: If usage is high and engagement is strong, prioritize Option A.

Output (this result): Recommended Option A

Why this one result happened

Global and Local Strategies

17

Textbook Review

Southwest Tech

LIME Recognition

Local Interpretable Model-Agnostic Explanations (LIME)

It is a local explainability technique.

For this course:

- ▶ recognize the idea
- ▶ do not implement it
- ▶ connect it to explaining one result

18

Textbook Review 

Ethics And Algorithms

Ethical questions include:

- ▶ Bias (& Typicality Bias)
- ▶ Discrimination
- ▶ Privacy
- ▶ Unfair outcomes
- ▶ Misleading evidence
- ▶ Responsibility for tool use



19

Textbook Review 

When Should We Use Algorithms?

Consider:

- ▶ cost
- ▶ time
- ▶ accuracy
- ▶ risk
- ▶ explainability
- ▶ human impact

Not every problem needs an algorithmic answer.



20





Explainability

21



Working Code Is Not The Whole Answer

Working code is valuable evidence.

But we still ask:

- What problem did it solve?
- What cases were tested?
- What did it assume?
- What could fail?
- What tradeoff did it make?



22



Global Explanation

Global explanation describes the overall approach.

Example:

"This solution checks eligibility by evaluating account status, training, equipment availability, and overdue-device status before returning a decision."



23



Local Explanation

Local explanation describes one output.

Example:

"This request returned 'needs review' because training was complete, but the device type requires supervisor approval."



24

Evidence Supports Explanation

Useful evidence may include:

- ▶ test table
- ▶ trace table
- ▶ timing table
- ▶ comparison table
- ▶ ranking table
- ▶ README explanation




25

Traceability


Traceability means a reviewer can follow:

- ▶ input
- ▶ logic
- ▶ evidence
- ▶ output
- ▶ explanation




26

Traceability: From Input to Explanation



Traceability connects what happened to how we explain it.

Traceability



27

Inconclusive Evidence

Inconclusive evidence does not fully prove the claim.

Examples:

- ▶ too few tests
- ▶ missing edge cases
- ▶ tiny dataset
- ▶ only happy-path examples
- ▶ unsupported performance claims





28

Misguided Evidence

Misguided evidence can look convincing while measuring the wrong thing.

Ask:

- ▶ Does this evidence match the claim?
- ▶ Is the test realistic?
- ▶ Is the dataset representative?
- ▶ Is the conclusion too broad?





29

Limitations Are Not Failure

A limitation is an honest boundary. It can explain:

- ▶ Where the solution may fail
- ▶ What data is missing
- ▶ What assumption may not hold
- ▶ What should be improved later




30



Tradeoff Explanation

Tradeoffs may involve:

- correctness
- readability
- speed
- memory
- simplicity
- fairness
- explainability

31



Part 2 Is Explanation, Not A Gotcha

Part 2 of the Final, asks you to explain your own submitted work.

You may be asked about:

- ▶ logic flow
- ▶ evidence
- ▶ assumptions
- ▶ edge cases
- ▶ tradeoffs
- ▶ AI/tool use
- ▶ improvements

32



Ethics and Evidence Quality

33




Bias And Discrimination

Bias can shape algorithmic results.

Watch for:

- ▶ incomplete data
- ▶ skewed examples
- ▶ unfair criteria
- ▶ historical patterns
- ▶ hidden assumptions

34



How Hidden Data Skew Affects Decisions

Training Data
What the model learns from

- Missing or uneven data
- Some groups overrepresented
- Some groups underrepresented
- Important cases may be missing

Model / Decision Process
Weights evidence to make a decision

Decision Result
What the system outputs

Uneven data can produce uneven outcomes. Better data leads to fairer, more reliable, and more trustworthy decisions.

Image Slide

35



Privacy

Privacy asks:

- ▶ What data is collected?
- ▶ Is it necessary?
- ▶ Who can see it?
- ▶ How is it protected?
- ▶ Could the output reveal sensitive information?

36



Problems With Learning Algorithms

Learning algorithms may:

- ▶ learn from poor data
- ▶ amplify harmful patterns
- ▶ drift over time
- ▶ behave differently after deployment
- ▶ become difficult to explain

ALGORITHM

NEXT

37



Ethical Questions By Context



Ask different questions for:

- ▶ classification
- ▶ regression
- ▶ recommendation
- ▶ data mining

Same principle:
What could go wrong, and who is affected?

This Photo by Unknown Author is licensed under CC BY-SA

38



Factors Affecting Algorithmic Solutions

Results can be affected by:

- ▶ data quality
- ▶ feature choice
- ▶ assumptions
- ▶ model or algorithm choice
- ▶ evaluation method
- ▶ deployment context
- ▶ human interpretation

39



Unfair Outcomes

An unfair outcome may come from:

- ▶ unfair criteria
- ▶ missing context
- ▶ biased data
- ▶ poor thresholds
- ▶ over-automation
- ▶ weak review process

This Photo by Unknown Author is licensed under CC BY-SA

40



Traceability As Protection

Traceability helps answer:

- ▶ What happened?
- ▶ Why did it happen?
- ▶ What evidence supports it?
- ▶ Who reviewed it?
- ▶ What should change?

41



AI Use Accountability


If AI helped, explain:

- ▶ What it helped with
- ▶ What you changed
- ▶ What you tested
- ▶ What you rejected
- ▶ What you still understand and own


42



43



44




45



46



47



48

Cost, Time, Accuracy



When deciding whether to use an algorithm, ask:

- What does it cost?
- How much time does it save?
- How accurate must it be?
- What happens when it is wrong?

49

Accuracy Is Not Enough



A solution may be accurate but still problematic if it is:

- ▶ impossible to explain
- ▶ unfair
- ▶ too expensive
- ▶ too slow
- ▶ too brittle
- ▶ harmful when wrong

50

Black Swan Events



Black swan events are:

- ▶ unexpected
- ▶ high impact
- ▶ easier to explain after they happen
- ▶ not always surprising to everyone

They expose prediction limits.

51

Forecasting Dilemmas



Hard questions:

- ▶ Can this be predicted?
- ▶ Can the implications be predicted?
- ▶ What data would be needed?
- ▶ What uncertainty remains?
- ▶ Who must decide anyway?

52

Post-Event Predictability



After an event, people may say:


"It was obvious."

But before the event:

- ▶ evidence may be incomplete
- ▶ signals may be noisy
- ▶ implications may be unclear

53

Practical Application: COVID-19



COVID-19 illustrates practical limits:

- ▶ changing conditions
- ▶ incomplete data
- ▶ uncertain implications
- ▶ high-stakes decisions
- ▶ need for revision over time

54



Be Careful With Claims

Strong final explanations avoid overclaiming.

Say:

- ▶ "My evidence shows..."
- ▶ "This assumes..."
- ▶ "This may fail when..."
- ▶ "A limitation is..."
- ▶ "With more time, I would..."

55



Southwest Tech

Final Bridge & Optional Bonus

56



Southwest Tech

Chapter 16 As Final Support

Chapter 16 can help you strengthen final explanations.

Use it for:

- ▶ explainability
- ▶ assumptions
- ▶ limitations
- ▶ bias awareness
- ▶ privacy awareness
- ▶ evidence quality
- ▶ tradeoff reasoning

57



Southwest Tech

Optional Enrichment / Bonus Credit

You may be able to earn bonus credit by clearly applying a Chapter 16 concept to your final work. (Instructors Discretion!!)


Examples:

- ▶ bias risk
- ▶ privacy risk
- ▶ inconclusive evidence
- ▶ black swan limitation
- ▶ cost/time/accuracy tradeoff

58

Southwest Tech

Optional Bonus Credit as an Evidence-Based Add-On



Final Submission

- ✓ Required components completed
- ✓ Meets assignment criteria
- ✓ Ready for evaluation

Optional Bonus
Explainability or ethics connection

example justification

Bonus requires clear evidence and explanation.

Optional Enrichment / Bonus Credit

59



Bonus Credit Must Be Demonstrated

Bonus-worthy application should:

- ▶ connect to your actual solution
- ▶ use evidence
- ▶ explain the risk or tradeoff
- ▶ avoid vague name-dropping
- ▶ appear in README/Part 2 explanation

60



FINAL Exam Structure



The final has two parts:


- ▶ Part 1 - Applied Solution Set
- ▶ Part 2 - Explanation Defense

Part 1 is what you submit.
Part 2 is how you explain it

61



Part 1 Readiness



Part 1 should include:

- ▶ working solutions
- ▶ tests and evidence
- ▶ assumptions
- ▶ limitations
- ▶ comparison where required
- ▶ AI/reference-use note when applicable

62



Part 2 Readiness



Be ready to explain:

- ▶ what your code does
- ▶ why a test passed or failed
- ▶ why you chose an approach
- ▶ what assumption matters
- ▶ what limitation remains
- ▶ what AI or previous work contributed

63



README As Explanation Prep



Your README helps you prepare Part 2.


It should organize:

- ▶ task summaries
- ▶ evidence tables
- ▶ assumptions
- ▶ limitations
- ▶ tradeoffs
- ▶ AI/tool-use notes
- ▶ previous-work notes

64



A Professional README Document, Prepared with Evidence



Documentation connects code, evidence, and explanation.

README Prep

65



Final Explanation Pattern



Use this pattern:

1. What did I build?
2. What evidence shows it works?
3. What did I assume?
4. What tradeoff did I make?
5. What limitation remains?
6. What did I verify?


66



67



68



69



70